

Cleaning and preparing the data for model training

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

PROBLEM STATEMENT

A retail company "ABC Private Limited" wants to understand the customer purchase behaviour (specifically, purchase amount) against various products of different categories. They have shared purchase summary of various customers for selected high volume products from last month. The data set also contains customer demographics (age,gender,marital status, city_type,stay_in_current_city), product details (product_id and product_category) and Total purchase_amount from last month.

Now, they want to build a model to predict and purchase amount of customer against various products which will help them to create personalized offer for customers against different products.

```
In [3]: # importing the dataset
df_train = pd.read_csv('blackfriday_train.csv')
```

```
In [4]: df_train.head(10)
```

```
Out[4]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Categr
0	1000001	P00069042	F	0-17	10	A	2	0	3	
1	1000001	P00248942	F	0-17	10	A	2	0	1	
2	1000001	P00087842	F	0-17	10	A	2	0	12	
3	1000001	P00085442	F	0-17	10	A	2	0	12	
4	1000002	P00285442	M	55+	16	C	4+	0	8	
5	1000003	P00193542	M	26-35	15	A	3	0	1	
6	1000004	P00184942	M	46-50	7	B	2	1	1	
7	1000004	P00346142	M	46-50	7	B	2	1	1	
8	1000004	P0097242	M	46-50	7	B	2	1	1	
9	1000005	P00274942	M	26-35	20	A	1	1	8	

```
In [5]: df_train.shape
```

```
Out[5]: (550068, 12)
```

In [6]: df_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                       550068 non-null  int64
8   Product_Category_1                  550068 non-null  int64
9   Product_Category_2                  376430 non-null  float64
10  Product_Category_3                  166821 non-null  float64
11  Purchase                             550068 non-null  int64
dtypes: float64(2), int64(5), object(5)
memory usage: 50.4+ MB
```

In [7]: df_train.describe()

Out[7]:

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
count	5.500680e+05	550068.000000	550068.000000	550068.000000	376430.000000	166821.000000	550068.000000
mean	1.003029e+06	8.076707	0.409653	5.404270	9.842329	12.668243	9263.968713
std	1.727592e+03	6.522660	0.491770	3.936211	5.086590	4.125338	5023.065394
min	1.000001e+06	0.000000	0.000000	1.000000	2.000000	3.000000	12.000000
25%	1.001516e+06	2.000000	0.000000	1.000000	5.000000	9.000000	5823.000000
50%	1.003077e+06	7.000000	0.000000	5.000000	9.000000	14.000000	8047.000000
75%	1.004478e+06	14.000000	1.000000	8.000000	15.000000	16.000000	12054.000000
max	1.006040e+06	20.000000	1.000000	20.000000	18.000000	18.000000	23961.000000

```
In [8]: df_train.isnull().sum()
```

```
Out[8]: User_ID          0
Product_ID          0
Gender              0
Age                0
Occupation          0
City_Category       0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category_1   0
Product_Category_2  173638
Product_Category_3  383247
Purchase            0
dtype: int64
```

```
In [9]: ## Import the test data
df_test=pd.read_csv("blackfriday_test.csv")
```

```
In [10]: df_test.head()
```

```
Out[10]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Categr
0	1000004	P00128942	M	46-50	7	B	2	1	1	
1	1000009	P00113442	M	26-35	17	C	0	0	3	
2	1000010	P00288442	F	36-45	1	B	4+	1	5	
3	1000010	P00145342	F	36-45	1	B	4+	1	4	
4	1000011	P00053842	F	26-35	1	C	1	0	4	

```
In [11]: df_test.shape
```

```
Out[11]: (233599, 11)
```

```
In [12]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 233599 entries, 0 to 233598
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               233599 non-null int64
1   Product_ID                            233599 non-null object
2   Gender                                233599 non-null object
3   Age                                    233599 non-null object
4   Occupation                             233599 non-null int64
5   City_Category                          233599 non-null object
6   Stay_In_Current_City_Years            233599 non-null object
7   Marital_Status                         233599 non-null int64
8   Product_Category_1                     233599 non-null int64
9   Product_Category_2                     161255 non-null float64
10  Product_Category_3                     71037 non-null  float64
dtypes: float64(2), int64(4), object(5)
memory usage: 19.6+ MB
```

```
In [13]: df_test.describe()
```

```
Out[13]:
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
count	2.335990e+05	233599.000000	233599.000000	233599.000000	161255.000000	71037.000000
mean	1.003029e+06	8.085407	0.410070	5.276542	9.849586	12.669454
std	1.726505e+03	6.521146	0.491847	3.736380	5.094943	4.125944
min	1.000001e+06	0.000000	0.000000	1.000000	2.000000	3.000000
25%	1.001527e+06	2.000000	0.000000	1.000000	5.000000	9.000000
50%	1.003070e+06	7.000000	0.000000	5.000000	9.000000	14.000000
75%	1.004477e+06	14.000000	1.000000	8.000000	15.000000	16.000000
max	1.006040e+06	20.000000	1.000000	18.000000	18.000000	18.000000

```
In [14]: df_test.isnull().sum()
```

```
Out[14]: User_ID          0
Product_ID          0
Gender              0
Age                0
Occupation          0
City_Category       0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category_1  0
Product_Category_2  72344
Product_Category_3  162562
dtype: int64
```

```
In [15]: ## Merging the train and test dataset
```

```
df = df_train.append(df_test)
```

C:\Users\harsh\AppData\Local\Temp\ipykernel_16220\2182507668.py:3: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
df = df_train.append(df_test)
```

In [16]: `df.head()`

Out[16]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Categr
0	1000001	P00069042	F	0-17	10	A	2	0	3	
1	1000001	P00248942	F	0-17	10	A	2	0	1	
2	1000001	P00087842	F	0-17	10	A	2	0	12	
3	1000001	P00085442	F	0-17	10	A	2	0	12	
4	1000002	P00285442	M	55+	16	C	4+	0	8	

In [17]: `# Dropping USER_ID column because it is of no use`

`df.drop(["User_ID"],axis=1,inplace=True)`

In [18]: `df.head()`

Out[18]:

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Pro
0	P00069042	F	0-17	10	A	2	0	3		NaN
1	P00248942	F	0-17	10	A	2	0	1		6.0
2	P00087842	F	0-17	10	A	2	0	12		NaN
3	P00085442	F	0-17	10	A	2	0	12		14.0
4	P00285442	M	55+	16	C	4+	0	8		NaN

Fixing categorical column into Numerical

In [19]: `# GENDER Column`

```
df["Gender"] = df["Gender"].map({"F":0, "M":1})
```

In [20]: `df.head()`

Out[20]:

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Pro
0	P00069042	0	0-17	10	A	2	0	3	NaN	
1	P00248942	0	0-17	10	A	2	0	1	6.0	
2	P00087842	0	0-17	10	A	2	0	12	NaN	
3	P00085442	0	0-17	10	A	2	0	12	14.0	
4	P00285442	1	55+	16	C	4+	0	8	NaN	

In [21]: `# AGE column`

```
df["Age"].unique()
```

Out[21]: `array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
 dtype=object)`

In [22]: `df["Age"] = df["Age"].map({'0-17':1, '18-25':2, '26-35':3, '36-45':4, '46-50':5, '51-55':6, '55+':7})`

In [23]: `df.head()`

Out[23]:

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Pro
0	P00069042	0	1	10	A	2	0	3	NaN	
1	P00248942	0	1	10	A	2	0	1	6.0	
2	P00087842	0	1	10	A	2	0	12	NaN	
3	P00085442	0	1	10	A	2	0	12	14.0	
4	P00285442	1	7	16	C	4+	0	8	NaN	

```
### 2nd technique:----->>>>> Using LABEL ENCODING
from sklearn import preprocessing

# Label_encoder object knows how to understand word label
label_encoder = preprocessing.LabelEncoder()

# Encode labels in column "Age"
df["Age"] = label_encoder.fit_transform(df["Age"])

df["Age"].unique()
```

In [24]: `# CITY_CATEGORY column`

```
df_city = pd.get_dummies(df['City_Category'], drop_first=True)
```

```
In [25]: df_city.head()
```

Out[25]:

	B	C
0	0	0
1	0	0
2	0	0
3	0	0
4	0	1

```
In [55]: df = pd.concat([df,df_city],axis=1)
df.head()
```

Out[55]:

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
0	P00069042	0	1	10	2	0	3	8.0	16.0
1	P00248942	0	1	10	2	0	1	6.0	14.0
2	P00087842	0	1	10	2	0	12	8.0	16.0
3	P00085442	0	1	10	2	0	12	14.0	16.0
4	P00285442	1	7	16	4	0	8	8.0	16.0

```
In [56]: ##### now dropping city_category  
  
df.drop('City_Category',axis=1,inplace=True)  
  
## already DROPPED this columnn thats why showing error
```

KeyError

Traceback (most recent call last)

Cell In[56], line 3

```

1 ##### now dropping city_category
----> 3 df.drop('City_Category',axis=1,inplace=True)

```

File ~\anaconda3\lib\site-packages\pandas\util_decorators.py:331, in deprecate_nonkeyword_arguments.<locals>.decorate.<locals>.wrapper(*args, **kwargs)

```

325 if len(args) > num_allow_args:
326     warnings.warn(
327         msg.format(arguments=_format_argument_list(allow_args)),
328         FutureWarning,
329         stacklevel=find_stack_level(),
330     )
--> 331 return func(*args, **kwargs)

```

File ~\anaconda3\lib\site-packages\pandas\core\frame.py:5399, in DataFrame.drop(self, labels, axis, index, columns, level, inplace, errors)

```

5251 @deprecate_nonkeyword_arguments(version=None, allowed_args=["self", "labels"])
5252 def drop( # type: ignore[override]
5253     self,
5254     (...)
5260     errors: IgnoreRaise = "raise",
5261 ) -> DataFrame | None:
5262     """
5263     Drop specified labels from rows or columns.
5264
5265     (...)
5397         weight 1.0      0.8
5398     """
-> 5399     return super().drop(
5400         labels=labels,
5401         axis=axis,
5402         index=index,
5403         columns=columns,
5404         level=level,
5405         inplace=inplace,
5406         errors=errors,
5407     )

```

File ~\anaconda3\lib\site-packages\pandas\util_decorators.py:331, in deprecate_nonkeyword_arguments.<locals>.decorate

```
te.<locals>.wrapper(*args, **kwargs)
    325 if len(args) > num_allow_args:
    326     warnings.warn(
    327         msg.format(arguments=_format_argument_list(allow_args)),
    328         FutureWarning,
    329         stacklevel=find_stack_level(),
    330     )
--> 331 return func(*args, **kwargs)
```

File ~\anaconda3\lib\site-packages\pandas\core\generic.py:4505, in NDFrame.drop(self, labels, axis, index, columns, level, inplace, errors)

```
    4503 for axis, labels in axes.items():
    4504     if labels is not None:
-> 4505         obj = obj._drop_axis(labels, axis, level=level, errors=errors)
    4507 if inplace:
    4508     self._update_inplace(obj)
```

File ~\anaconda3\lib\site-packages\pandas\core\generic.py:4575, in NDFrame._drop_axis(self, labels, axis, level, errors, only_slice)

```
    4573 labels_missing = (axis.get_indexer_for(labels) == -1).any()
    4574 if errors == "raise" and labels_missing:
-> 4575     raise KeyError(f"{labels} not found in axis")
    4577 if is_extension_array_dtype(mask.dtype):
    4578     # GH#45860
    4579     mask = mask.to_numpy(dtype=bool)
```

KeyError: "['City_Category'] not found in axis"

```
In [ ]: df.head()
```

MISSING VALUES

```
In [29]: df.isnull().sum()
```

```
Out[29]: Product_ID          0
Gender          0
Age            0
Occupation      0
Stay_In_Current_City_Years  0
Marital_Status  0
Product_Category_1  0
Product_Category_2  245982
Product_Category_3  545809
Purchase        233599
B              0
C              0
dtype: int64
```

we now focus on replacing missing values

```
In [30]: ### replacing the "missing values" using mode for "PRODUCT_CATEGORY_2"
```

```
In [31]: df["Product_Category_2"].unique()
```

```
Out[31]: array([nan,  6., 14.,  2.,  8., 15., 16., 11.,  5.,  3.,  4., 12.,  9.,
        10., 17., 13.,  7., 18.])
```

```
In [32]: df['Product_Category_2'].value_counts()
```

```
Out[32]: 8.0      91317
        14.0     78834
        2.0     70498
        16.0     61687
        15.0     54114
        5.0     37165
        4.0     36705
        6.0     23575
        11.0    20230
        17.0    19104
        13.0    15054
        9.0     8177
        12.0     7801
        10.0     4420
        3.0     4123
        18.0     4027
        7.0      854
        Name: Product_Category_2, dtype: int64
```

```
In [33]: # filling the "missing values" using mode for "PRODUCT_CATEGORY_2"
```

```
df["Product_Category_2"] = df["Product_Category_2"] \
    .fillna(df['Product_Category_2'].mode()[0])
```

```
In [34]: df["Product_Category_2"].isnull().sum()
```

```
Out[34]: 0
```

```
In [35]: # replacing the "missing values" using mode for "PRODUCT_CATEGORY_3"
```

```
In [36]: df["Product_Category_3"].unique()
```

```
Out[36]: array([nan, 14., 17.,  5.,  4., 16., 15.,  8.,  9., 13.,  6., 12.,  3.,
                18., 11., 10.])
```

```
In [37]: df["Product_Category_3"].value_counts()
```

```
Out[37]: 16.0    46469
          15.0    39968
          14.0    26283
          17.0    23818
           5.0    23799
           8.0    17861
           9.0    16532
          12.0    13115
          13.0     7849
           6.0     6888
          18.0     6621
           4.0     2691
          11.0     2585
          10.0     2501
           3.0       878
          Name: Product_Category_3, dtype: int64
```

```
In [38]: # filling the "missing values" using mode for "PRODUCT_CATEGORY_3"
```

```
df["Product_Category_3"] = df["Product_Category_3"].fillna(df["Product_Category_3"].mode()[0])
```

```
In [39]: df["Product_Category_3"].isnull().sum()
```

```
Out[39]: 0
```


In [40]: `df.head()`

Out[40]:

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
0	P00069042	0	1	10	2	0	3	8.0	16.0
1	P00248942	0	1	10	2	0	1	6.0	14.0
2	P00087842	0	1	10	2	0	12	8.0	16.0
3	P00085442	0	1	10	2	0	12	14.0	16.0
4	P00285442	1	7	16	4+	0	8	8.0	16.0

In [41]: `## Removing "+" in "Stay_In_Current_City_Years"`

```
df["Stay_In_Current_City_Years"].unique()
```

Out[41]: `array(['2', '4+', '3', '1', '0'], dtype=object)`

In [42]: `df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].str.replace('+','')`

C:\Users\harsh\AppData\Local\Temp\ipykernel_16220\2063355665.py:1: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].str.replace('+','')
```

In [43]: `df.head()`

Out[43]:

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
0	P00069042	0	1	10	2	0	3	8.0	16.0
1	P00248942	0	1	10	2	0	1	6.0	14.0
2	P00087842	0	1	10	2	0	12	8.0	16.0
3	P00085442	0	1	10	2	0	12	14.0	16.0
4	P00285442	1	7	16	4	0	8	8.0	16.0

CONVERTING Dtype

In [44]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Product_ID                            783667 non-null object
1   Gender                                783667 non-null int64
2   Age                                    783667 non-null int64
3   Occupation                            783667 non-null int64
4   Stay_In_Current_City_Years            783667 non-null object
5   Marital_Status                        783667 non-null int64
6   Product_Category_1                    783667 non-null int64
7   Product_Category_2                    783667 non-null float64
8   Product_Category_3                    783667 non-null float64
9   Purchase                              550068 non-null float64
10  B                                      783667 non-null uint8
11  C                                      783667 non-null uint8
dtypes: float64(3), int64(5), object(2), uint8(2)
memory usage: 67.3+ MB
```

In [45]: *## Converting Dtype from 'Object' to "int64" of Column "Stay_In_Current_City_Years"*

```
df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].astype(int)

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   Product_ID                           783667 non-null object  
 1   Gender                               783667 non-null int64   
 2   Age                                  783667 non-null int64   
 3   Occupation                           783667 non-null int64   
 4   Stay_In_Current_City_Years          783667 non-null int32   
 5   Marital_Status                       783667 non-null int64   
 6   Product_Category_1                   783667 non-null int64   
 7   Product_Category_2                   783667 non-null float64  
 8   Product_Category_3                   783667 non-null float64  
 9   Purchase                             550068 non-null float64  
10   B                                     783667 non-null uint8   
11   C                                     783667 non-null uint8   
dtypes: float64(3), int32(1), int64(5), object(1), uint8(2)
memory usage: 64.3+ MB
```

In [46]: *## ## Converting Dtype from 'uint8' to "int" of Column "B & C"*

```
df['B']=df['B'].astype(int)
df['C']=df['C'].astype(int)
```

```
In [47]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Product_ID                           783667 non-null object
 1   Gender                               783667 non-null int64
 2   Age                                  783667 non-null int64
 3   Occupation                           783667 non-null int64
 4   Stay_In_Current_City_Years          783667 non-null int32
 5   Marital_Status                       783667 non-null int64
 6   Product_Category_1                  783667 non-null int64
 7   Product_Category_2                  783667 non-null float64
 8   Product_Category_3                  783667 non-null float64
 9   Purchase                             550068 non-null float64
10   B                                    783667 non-null int32
11   C                                    783667 non-null int32
dtypes: float64(3), int32(3), int64(5), object(1)
memory usage: 68.8+ MB
```

VISUALIZATION

In [58]: `## Age Vs Purchase`

```
sns.barplot(df, x='Age', y='Purchase', hue='Gender')
```

Out[58]: `<Axes: xlabel='Age', ylabel='Purchase'>`




Inference: Purchasing of men is high then women

In [63]: *## Purchase with Occupation*

```
sns.barplot(df,x='Occupation',y='Purchase',hue='Gender')
```


Out[63]: <Axes: xlabel='Occupation', ylabel='Purchase'>



In [64]: `## Product_Category_1 with Purchase`

```
sns.barplot(df,x='Product_Category_1',y='Purchase',hue='Gender')
```


Out[64]: `<Axes: xlabel='Product_Category_1', ylabel='Purchase'>`



In [65]: `## Product_Category_2 with Purchase`

```
sns.barplot(df,x='Product_Category_2',y='Purchase',hue='Gender')
```


Out[65]: `<Axes: xlabel='Product_Category_2', ylabel='Purchase'>`



In [67]: *## Product_Category_3 with Purchase*

```
sns.barplot(data=df,x='Product_Category_3',y='Purchase',hue='Gender')
```

Out[67]: <Axes: xlabel='Product_Category_3', ylabel='Purchase'>



In []: *## Purchase with null values*

```
df_test=df[df['Purchase'].isnull()]\ndf_test
```

FEATURE SCALING

In []: *## Purchase without null values*

```
df_train = df[~df['Purchase'].isnull()]  
df_train
```

In []: *### feature scaling*

```
from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()
```

In []: df

In []:

In []:

In []: